

GraphBLAS API in Functional Style

1st Given Name Surname
dept. name of organization (of Aff.)
name of organization (of Aff.)
City, Country
email address or ORCID

2nd Given Name Surname
dept. name of organization (of Aff.)
name of organization (of Aff.)
City, Country
email address or ORCID

3rd Given Name Surname
dept. name of organization (of Aff.)
name of organization (of Aff.)
City, Country
email address or ORCID

Semyon Grigorev
Saint Petersburg State University,
JetBrains Research,
St. Petersburg, Russia
s.v.grigoriev@spbu.ru,
semyon.grigorev@jetbrains.com

[illegible]

Index Terms—graph analysis, sparse linear algebra, Graph-BLAS API, GPGPU, parallel programming, functional programming, .NET, OpenCL

I. INTRODUCTION

One of the promising ways to high-performance graph analysis is based on utilization of linear algebra: operations over vectors and matrices can be efficiently implemented on modern parallel hardware, and once we reduce given graph analysis problem to composition of such operations, we get high-performance solution for our problem. Well-known example of such reduction is a reduction of all-pairs shortest path (APSP) problem to matrix multiplication over appropriate *semiring*. To formalize, generalize this observation and make it useful in practice, GraphBLAS API standard was found [1]. GraphBLAS API introduces appropriate algebraic structures (monoid, semiring), objects (scalar, vector, matrix), and operations over them such that to form building blocks to create graph analysis algorithms. Not only Graphs. Sparse Linear algebra.

Reference implementation SuiteSparse:GraphBLAS [2].
Other implementations: CombBLAS, !!! .

GPGPU for high-performance analysis of huge amount of data. GraphBLAST [3] — GraphBLAST in CUDA C.

High-level programming languages for application development vs low-level for high-performance programming. Moreover, specific languages for GPGPU programming: CUDA C,

OpenCL C. Problems with types, compositionality, optimizations (kernel fusion).

Functional programming. Type systems. Optimizations. Futhark [6], kernel fusion, specialization, deforestation etc.

In this work we discuss some !!!, provide !!! which we start working on. It is implemented on .NET platform in F# programming language with OpenCL backend. Portability of OpenCL: CPU, GPGPU, FPGA [4], [5]. A way to hardware acceleration of sparse linear algebra and graph analysis. **Brahma.FSharp**¹. Our preliminary evaluation shows that !!!!

II. DESIGN PRINCIPLES

Functional style, types, optimizations, etc.

Code example with description and explanations.

III. IMPLEMENTATION DETAILS

Details on implementation.

Architecture.

IV. EVALUATION

Evaluation of the proposed implementation.

Hardware configuration description.

SuiteSparse, Math.NET Numerics², GraphBLAST, ???, and our solution on CPU and GPGPU.

Dataset description.

Results.

Results analysis. and conclusion.

V. CONCLUSION

Conclusion, current state, results.

Future work. Library extension up to full GraphBLAS API implementation.

Evaluation. Comparison with other implementations. Manual implementation versus translation.

Brahma.FSharp improvements. Algebraic data types. Example with Min-Plus or other semiring.

1 !!!

²Library which provides numerical computations primitives for .NET: <https://numerics.mathdotnet.com/>. Access date: 12.01.2021.

Identify applicable funding agency here. If none, delete this.

REFERENCES

- [1] J. Kepner, P. Aaltonen, D. Bader, A. Buluc, F. Franchetti, J. Gilbert, D. Hutchison, M. Kumar, A. Lumsdaine, H. Meyerhenke, S. McMillan, C. Yang, J. D. Owens, M. Zalewski, T. Mattson, and J. Moreira, “Mathematical foundations of the graphblas,” in *2016 IEEE High Performance Extreme Computing Conference (HPEC)*, Sep. 2016, pp. 1–9.
- [2] T. A. Davis, “Algorithm 1000: Suitesparse:graphblas: Graph algorithms in the language of sparse linear algebra,” *ACM Trans. Math. Softw.*, vol. 45, no. 4, Dec. 2019. [Online]. Available: <https://doi.org/10.1145/3322125>
- [3] C. Yang, A. Buluç, and J. D. Owens, “GraphBLAST: A high-performance linear algebra-based graph framework on the GPU,” *arXiv preprint*, 2019.
- [4] T. Kenter, “Invited tutorial: Opencl design flows for intel and xilinx fpgas: Using common design patterns and dealing with vendor-specific differences,” in *FSP Workshop 2019; Sixth International Workshop on FPGAs for Software Programmers*. VDE, 2019, pp. 1–8.
- [5] K. Shagrirhaya, K. Kepa, and P. Athanas, “Enabling development of opencl applications on fpga platforms,” in *2013 IEEE 24th International Conference on Application-Specific Systems, Architectures and Processors*, 2013, pp. 26–30.
- [6] T. Henriksen, N. G. W. Serup, M. Elsmann, F. Henglein, and C. E. Oancea, “Futhark: Purely functional gpu-programming with nested parallelism and in-place array updates,” in *Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI 2017. New York, NY, USA: ACM, 2017, pp. 556–571. [Online]. Available: <http://doi.acm.org/10.1145/3062341.3062354>